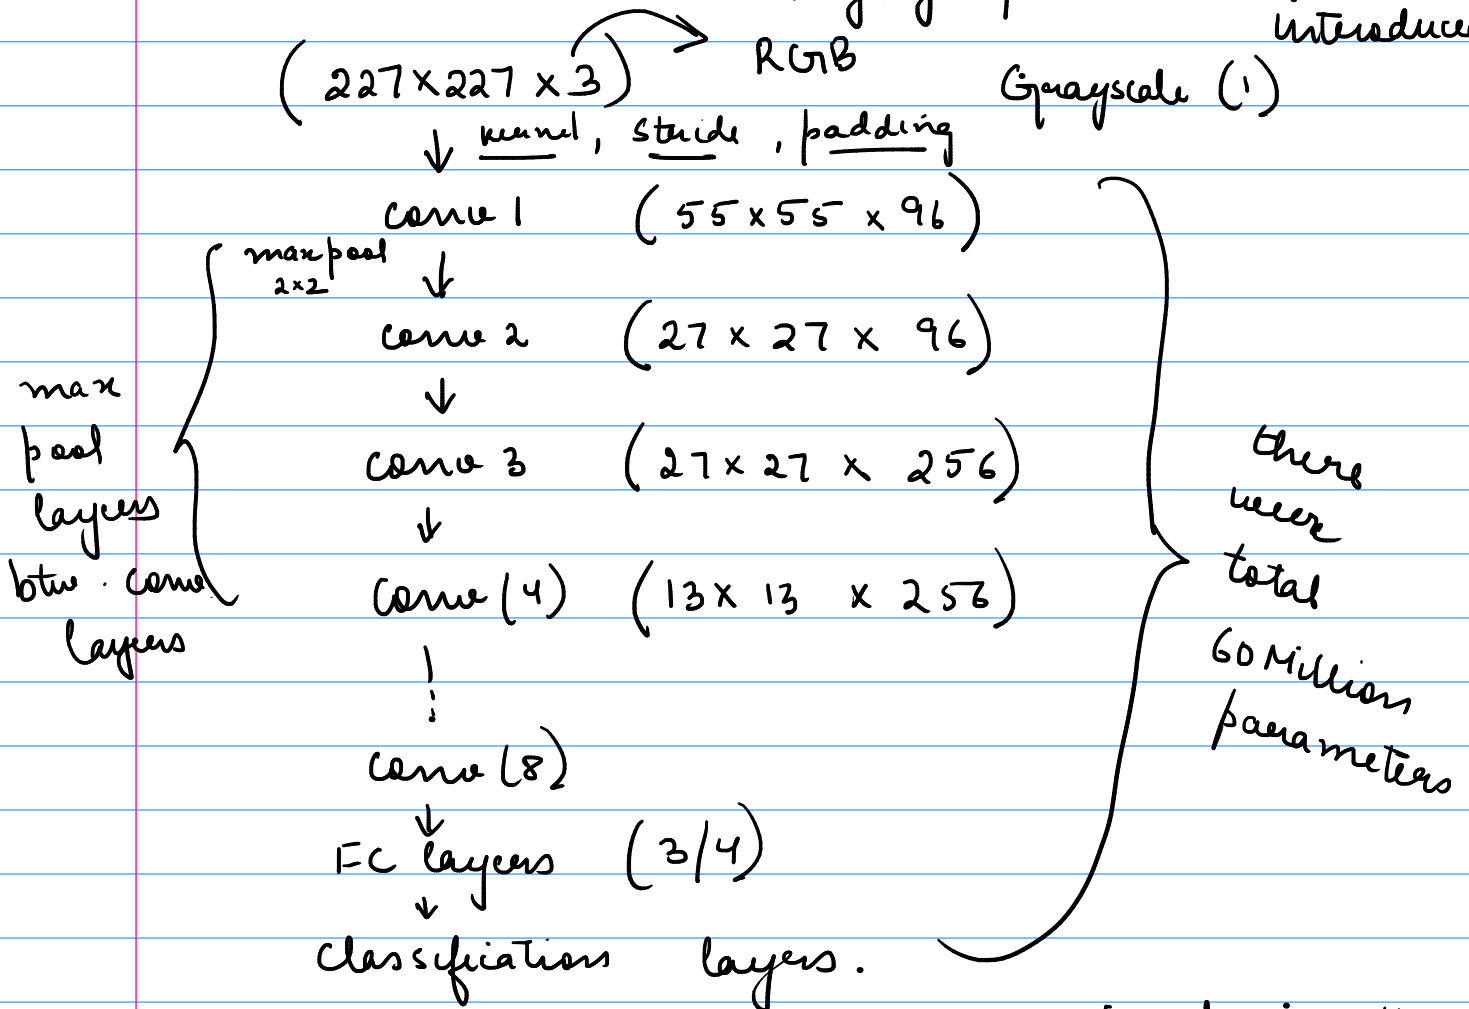



Week-4

- ILSVRC \rightarrow ImageNet Large Scale Visual Recognition Challenge.
 - \rightarrow 1000 categories \Rightarrow task of classification among these categories.
 - \rightarrow Metric \Rightarrow top-5 error rate %.

1. AlexNet (2012) \Rightarrow 28% \rightarrow 16% \Rightarrow CNNs were introduced



- \rightarrow used 2 GPUs
- \rightarrow use of ReLU Activation & dropout were introduced for the 1st time
 - \rightarrow only during the train time to make sure there is no bias.
 - \rightarrow $\max(0, a)$



kernel = 2×2

stride = 1

padding = 0

3×3

3×3


2×2

$\rightarrow \sim 8\%$

2. VGGNet (2014)

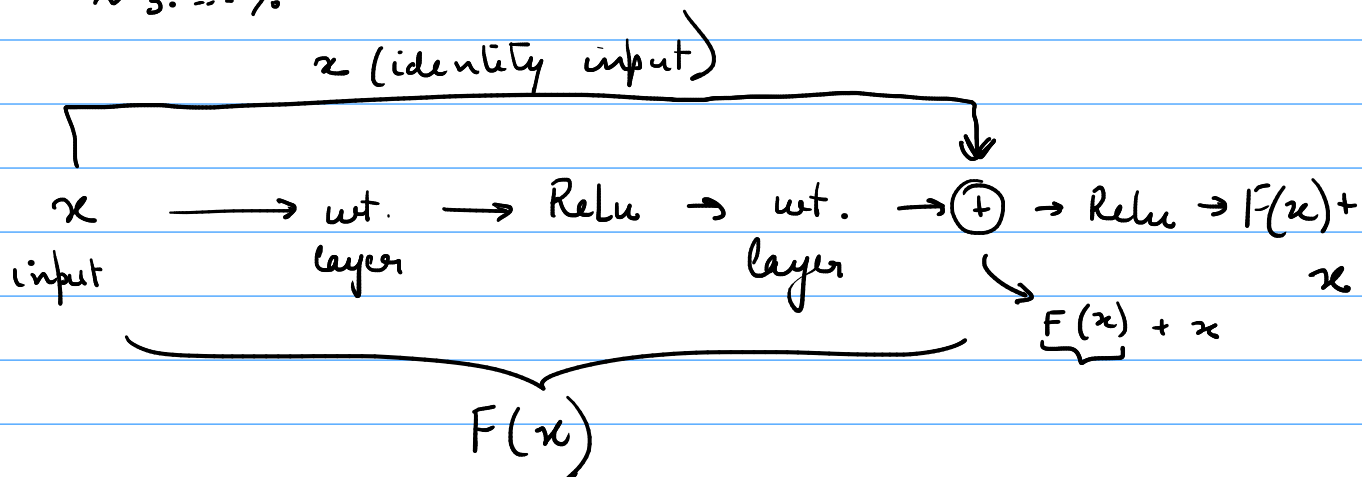
- \rightarrow 16-19 layers \Rightarrow extremely dense networks
- \rightarrow performance better
- \rightarrow in the networks only 3×3 ^{kernel} layers were used.
- \rightarrow they stacked up these 3×3 layers \hookrightarrow larger effective receptive field.

\Rightarrow Challenges \rightarrow when the network depth ^{no. of layers} started increasing beyond a certain pt., then accuracies tend to saturate & then decay exponentially.



3. ResNet \Rightarrow Residual block \Rightarrow introduces a skip connection.

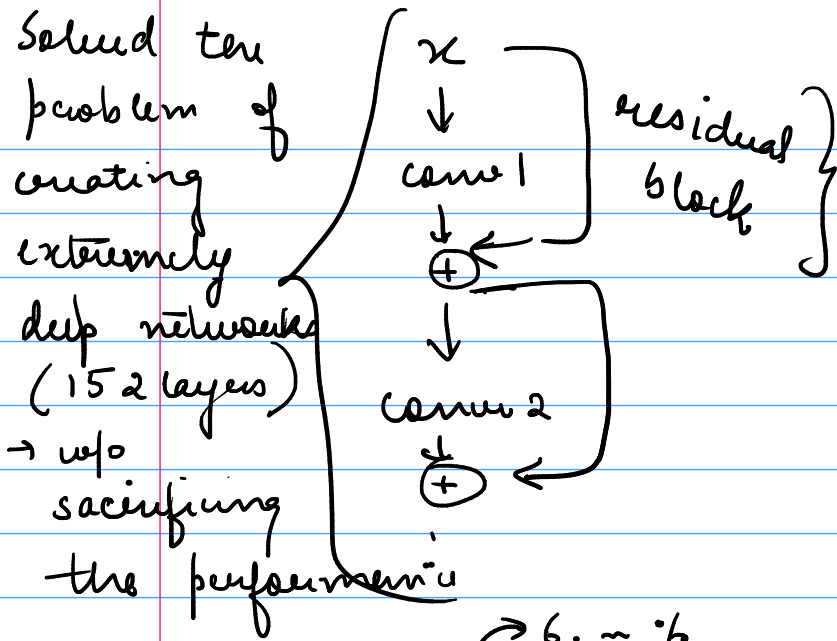
$\sim 3 \dots \%$




\rightarrow It adds block x input directly to the output $F(x) \Rightarrow F(x) + H(x) + x \Rightarrow$ Residual fun

$$A + B =$$

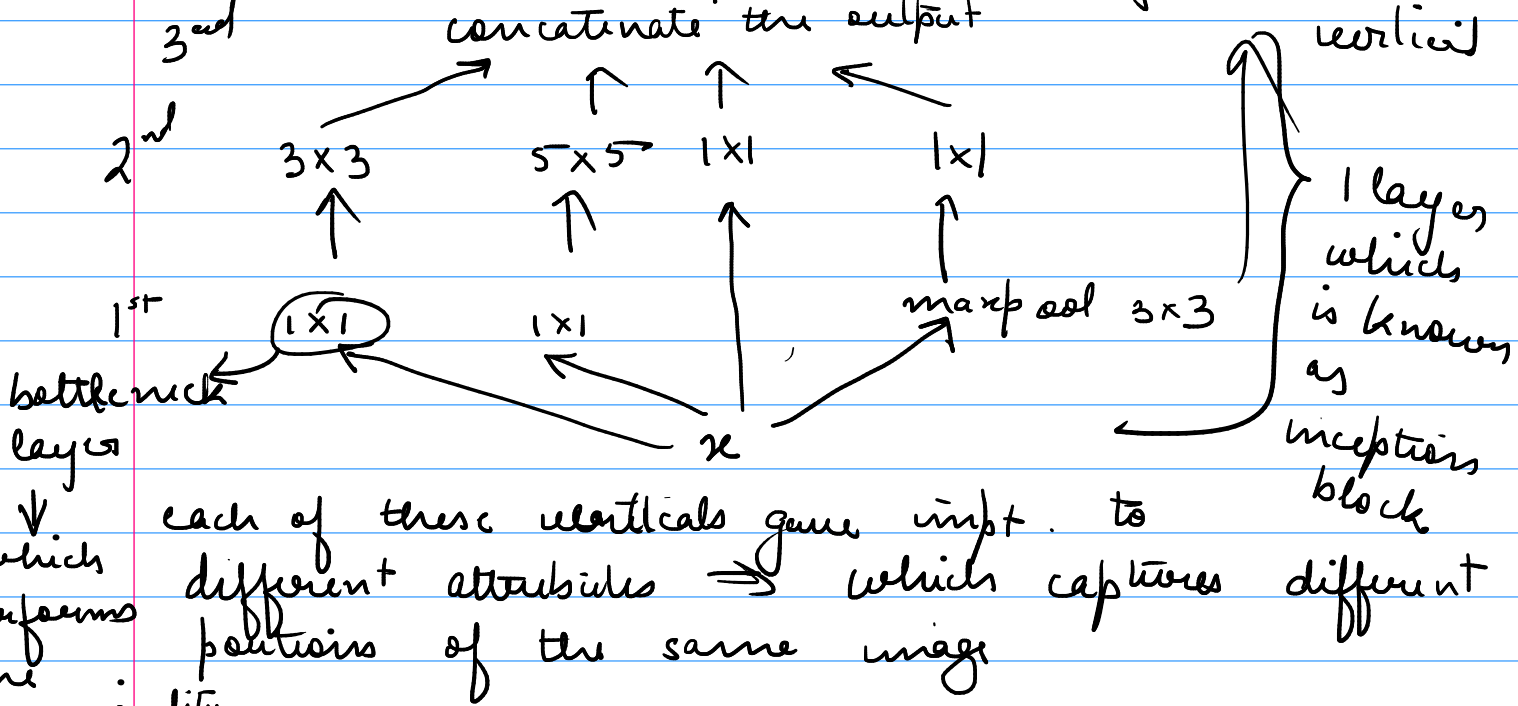
Solved the problem of creating extremely deep networks (152 layers) \rightarrow w/o sacrificing the performance



which brought the skip connection



4. GoogleNet \Rightarrow Inception module \rightarrow 6. ~ %
3rd \rightarrow concatenate the output \rightarrow send this to next layer

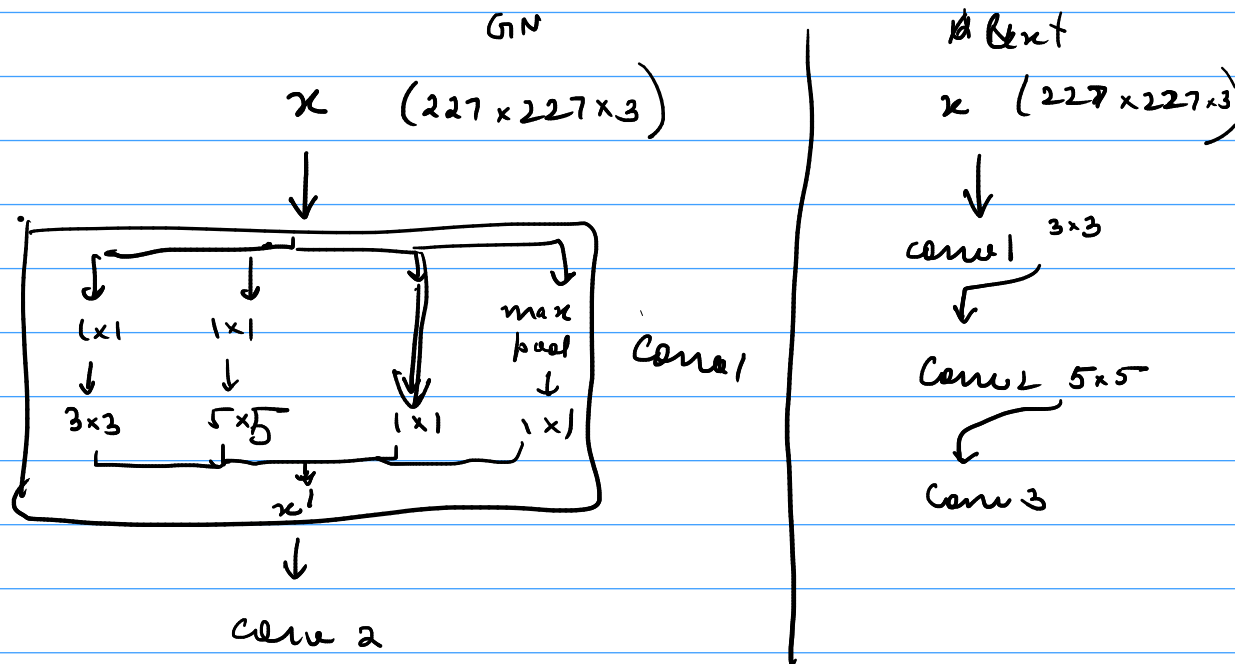


\rightarrow they converted the entire vertical networks into a cluster of both vertical & horizontal networks! \Rightarrow parallel conv. \rightarrow parallel convolution was performed by \rightarrow kernels $\rightarrow 1 \times 1, 3 \times 3, 5 \times 5$ \rightarrow max pool

\rightarrow even though, the no. of params \downarrow , it is computationally efficient

Adaptive Pool \longleftrightarrow PyTorch

→ FC layers was replaced GAP (global avg. pooling) w/o using the parameters



- Transfer learning
 ↳ how we utilize a pre-existing architecture to prepare for a specific downstream tasks.

General process -

1. use a pre-trained conv. base (nets & parameters)
2. we will remove one single classifier head (1000) layer.
3. add our own custom classifier head.
4. just train it.

- VGGNet → Freeze all its layers & replace the classifier head.
 ↳ only if your dataset is small

- AlexNet → unfreeze all the nets. & use the new data to retrain the model either partially or fully. ⇒ Fine-tuning